

MONDAY, SEPTEMBER 14TH, 2026

PHILOSOPHY

# The Sidewalk Was Empty

*Two drivers make the same choice on the same night, and everything that separates them happens after they have stopped deciding.*

3,149 words · 17-minute read

## The Arithmetic of Getting Home

Begin with a ratio that nobody set out to compute. The Centers for Disease Control, working from a telephone survey that asks adults whether they have driven in the past thirty days when they had perhaps too much to drink, puts alcohol-impaired driving episodes in the United States at roughly 125 million a year [1]. The two published analyses that show their working give 112 million episodes for 2010 and 121 million for 2012 [2], [3], and both pass a quiet internal test: dividing those totals by the reported rates of 479 and 505 episodes per thousand adults recovers adult populations of 233.8 and 239.6 million, which is the American adult population in those years to within a fraction of a per cent.

In 2023, 12,429 people died in crashes in which a driver was at or above the legal blood-alcohol limit — thirty per cent of all American road deaths, or one death every forty-two minutes [4]. Divide the deaths by the episodes and a rate falls out: one death for every 10,057 times somebody in America drives drunk. Because a single crash can kill more than one person, that ratio is an upper bound on the per-episode probability. Call it one in ten thousand and treat the figure as generous.

| How one episode ends        | Count in a year | Per episode | About one in |
|-----------------------------|-----------------|-------------|--------------|
| Nothing recorded            | 124.2 million   | 0.99346     | 1            |
| Arrest for impaired driving | 804,926         | 0.006439    | 155          |
| A death in the crash        | 12,429          | 0.000099    | 10,057       |

The three counts come from different years — the survey basis is 2020, the arrest estimate 2024 [5], the fatality count 2023 — and a self-report of “too much” is not the legal threshold used in the crash statistics. Neither caveat moves the orders of magnitude, which is the whole of the point. The overwhelmingly most likely outcome of driving drunk is that you get home, park badly, and go to bed.

Consider two people who leave the same bar at the same hour with the same alcohol in them and the same distance to drive. One arrives home. The other kills someone. The first is a private embarrassment; the second is a defendant, a headline, and for the rest of their life a person who killed somebody.

Nothing they did differs. Nothing they intended, chose, foresaw, or controlled differs. What differs is who else happened to be on the road at 1:14 in the morning.

Thomas Nagel put the case in a single sentence: if someone has had too much to drink and his car swerves onto the sidewalk, he can count himself morally lucky if there are no pedestrians in its path — and if there were, he would be to blame for their deaths and would probably be prosecuted for manslaughter [6]. The sidewalk is empty or it is not. The driver has no vote.

## A Symposium in July 1976

---

The name for this is moral luck, and it was fixed in the literature on a single afternoon. Bernard Williams and Nagel read papers to a joint session of the Aristotelian Society and the Mind Association in 1976; both were titled “Moral Luck,” and both were printed in the same supplementary volume, Williams at pages 115 to 135 and Nagel at 137 to 151 [6], [7]. Williams later said he had expected the phrase to sound like an oxymoron, and that the opposition it aroused was the point: a large part of what morality is for is to provide a shelter against luck [8].

The shelter is built out of a principle so obvious that it is rarely stated. Call it the control principle: people are morally assessable only for what is within their control. It is the reason we excuse the sleepwalker, the coerced, the person whose drink was spiked. Kant built an entire ethics on its most uncompromising form. Even if, owing to the special disfavour of fortune or the niggardly provision of a step-motherly nature, a good will should wholly lack power to accomplish its purpose, then, like a jewel, it would still shine by its own light, as a thing which has its whole value in itself [9]. Consequences are the weather. The will is the thing.

Nagel’s move was to notice that the control principle, applied honestly, does not stop where we would like it to. He sorted luck into four kinds. Resultant luck is luck in how things turn out: the swerve, the empty sidewalk. Circumstantial luck is luck in what situations you face at all — a man who left Germany for Argentina in 1930 on business led a quiet and harmless life, and never found out what he would have become had he stayed. Constitutive luck is luck in what you are like: your temperament, your patience, your capacity for cruelty, none of which you selected. Causal luck is luck in being the product of antecedent causes at all, which is the free will problem in a different hat [7]. Run the principle through all four and what survives is not a purified moral self but, in Nagel’s phrase, an extensionless point. That is the structure of the problem and the reason it has not been solved: the control principle is not a rule we are free to abandon, and it eats everything.

Williams came at it from the other side, with a case about which our judgements are not squeamish but enthusiastic. His Gauguin — a fiction attached to a real name — abandons his family for Tahiti and painting. Whether that choice was justified, Williams argued, cannot be settled in advance of knowing whether he became a great painter, and that turned on facts about his talent he could not have known when he left [6]. Success does not merely reveal that he was right; it is part of what makes him right. If that holds, retrospective luck is not an intruder in moral judgement. It is load-bearing.

## The Irregularity Adam Smith Refused to Fix

---

None of this was new in 1976. It was in print in 1759, in the least-read half of Adam Smith's *The Theory of Moral Sentiments*, whose Part II, Section III is titled "Of the Influence of Fortune upon the Sentiments of Mankind, with regard to the Merit or Demerit of Actions" [10]. Smith's description is exact. Fortune diminishes our sense of the merit or demerit of actions which arose from the most laudable or blameable intentions when they fail of producing their proposed effects, and increases it when they produce extraordinary consequences. The effects of an action, he wrote, often throw a shadow of merit or demerit upon the agent, though in his intention there was nothing that deserved either praise or blame. He called it, flatly, an irregularity in the sentiments of all men.

What Smith does next is what makes him worth reading here. Having diagnosed the irregularity, he declines to cure it, and gives it three chapters, the last titled "Of the Final Cause of this Irregularity of Sentiments." Nature, he argues, seems to have implanted it with the happiness and perfection of the species in view. Suppose we punished bad intentions as we punish bad outcomes. Then, in the line that has outlived everything around it, if the baseness of the thought which had given birth to no action seemed in the eyes of the world as much to call aloud for vengeance as the baseness of the action, every court of judicature would become a real inquisition [10].

So the irregularity is not a defect in moral cognition. It is a load-limiter. By attaching blame to visible outcomes rather than to invisible states of mind, a society buys itself a rule of evidence that no tribunal can abuse, at the cost of treating two identical people differently. Smith thought the trade worth making, and said so two hundred and seventeen years before the phrase "moral luck" existed. Whether he was right turns out to be answerable in a way he could not have managed, because his claim is about the evidential value of outcomes, and evidential value is something we now know how to compute.

## Smith's Defence, Written as Probability

---

Suppose agents come in two kinds. A careless one causes harm with probability  $p_1$  on any given occasion; a careful one with probability  $p_0$ . We do not observe which kind a person is; we observe only whether harm occurred, an event  $H$ . Bayes's rule gives the posterior odds that an agent was careless as the prior odds multiplied by a likelihood ratio:

$$\frac{P(C | H)}{P(\neg C | H)} = \frac{P(C)}{P(\neg C)} \cdot \frac{p_1}{p_0}$$

and, when no harm occurred, by the complementary ratio  $(1 - p_1)/(1 - p_0)$ .

The asymmetry between those two multipliers is the entire argument. Take a careless rate of one in a thousand and a careful rate of one in a hundred thousand — a hundred-fold difference in danger, which is not extravagant for drunk against sober driving. The harm multiplier is 100. The no-harm multiplier is

$0.999/0.99999 = 0.99901$ , which is to say one. Starting from a prior of 0.10 that a given driver was reckless, a death moves the posterior to 0.917. Getting home safely moves it to 0.09991. A rare harm is a shout; its absence is not even a whisper.

#### TWO QUESTIONS THAT SOUND LIKE ONE

Does this death tell us the driver was careless? Does this death tell us which of the careless drivers this one was? One of them has an excellent answer.

That is Smith's final cause, restated with a subscript. Outcome-tracking is a cheap, high-precision filter on an unobservable, and it requires no inquisitor to look inside anybody's head. A society that blamed on intentions alone would need an instrument for reading intentions, and the only such instrument is a tribunal empowered to decide what you were thinking. Smith's inquisition is not a rhetorical flourish; it is the alternative.

But notice what the same numbers say about coverage. Of all the careless agents in the population, the fraction who actually cause harm is  $p_1$  — one in a thousand. The filter has superb precision and essentially no recall. It convicts almost nobody it should not, and ignores 999 of every 1,000 people it should be worried about. In the driving numbers the gap is starker still: 124 million episodes a year end with nothing recorded at all.

## What the Same Numbers Cannot Do


Now ask the second question in the call-out, the one about which careless driver. Let each agent carry their own per-episode probability of causing harm, a number  $c$  that measures how dangerous they actually are, and let the outcome of a single episode be a draw from a coin of that bias. The law of total variance splits the variation in outcomes into the part that comes from real differences between agents and the part that is the coin:

$$\text{Var}(H) = \mathbb{E}[c(1-c)] + \text{Var}(c)$$

The first term is chance; the second is culpability. And because the correlation between an agent's true  $c$  and the observed outcome works out to the square root of the second term's share, the split can be read directly as how much of who gets blamed is really about who they are.

Put the American driving rate in. Let the average per-episode probability be the one in 10,057 computed above, and let drivers differ from one another substantially — a lognormal spread wide enough that the ninetieth-percentile driver is thirteen times as dangerous as the tenth. Then  $\text{Var}(c)$  is  $1.70 \times 10^{-8}$  against an  $\mathbb{E}[c(1-c)]$  of  $9.94 \times 10^{-5}$ , and the culpability share of the variance is 0.00017. The correlation between how dangerous a driver truly is and whether this particular drive killed anybody is 0.013. Forty million simulated agents reproduce that to three decimal places.

Widen the spread absurdly, to a hundred-and-sixty-eight-fold gap between the tenth and ninetieth percentiles, and the correlation climbs only to 0.073. The conclusion is robust because it does not depend on the spread: when harm is rare, almost all of the variation in who causes it is the coin, no matter how unequal the coins are.



**FIGURE** Left: the posterior probability that an agent was careless, given that harm did or did not occur, as a function of the prior, at a hundred-fold difference in harm rates. The no-harm curve is graphically indistinguishable from the line of no revision. Right: the correlation between an agent's true per-episode dangerousness and their observed harm record, against the number of episodes observed, for three widths of the underlying spread in dangerousness.

The same arithmetic answers a natural follow-up. If a single outcome tells you almost nothing about an individual, how many outcomes would? Observing  $n$  episodes multiplies the signal term by  $n$ , so the correlation reaches 0.5 — a merely respectable measurement — at about 1,950 episodes. At the observed rate of roughly twenty-nine impaired drives a year per person who does it at all, that is sixty-eight years of drinking and driving before the record of outcomes becomes a half-decent estimate of the driver.

So both things are true of the same number, and this is the part that does not fit in a slogan. Harm is powerful evidence that somebody in the vicinity was careless, and a nearly worthless measurement of which careless person you have got: excellent precision, no resolution. Michael Zimmerman reached the same shape of conclusion by argument rather than arithmetic — luck affects the scope of responsibility, what you are responsible for, without affecting its degree [11] — and Judith Jarvis Thomson made a parallel move, denying that the killer is more blameworthy while conceding that the world treats him differently [12]. Robert Hartman argues the other way, that luck really does change how blameworthy a person is [13]. What the numbers add to a live dispute is that if outcome-based blame is a way of measuring culpability, it is a bad one, and its defenders have to want it for something else.

## The Lottery in the Statute Book

---

The law runs the same machinery with the consequences visible. Nowhere is it starker than in the gap between attempt and completion. In England and Wales, murder carries a mandatory life sentence with a minimum term whose ordinary starting point for an adult is fifteen years; attempted murder has no mandatory minimum at all, and the definitive guideline's offence range opens at three years' custody [14]. Two people fire identical shots with identical intent. One hits.

David Lewis noticed that this is not merely an inconsistency but a specific institution, and gave it a name: the penal lottery [15]. We do not punish attempted murderers less because they are less wicked; almost everybody agrees they are not. What we do is issue them a ticket. The attempter and the murderer draw from the same urn, and the difference in their sentences is the difference between a winning and a losing ticket, drawn after the act rather than before. Lewis's provocation is that a penal lottery may be perfectly defensible if the risk of punishment is itself a punishment — we accept risk as a currency everywhere else — but that if we accept it here, we should say so out loud and stop pretending the murderer did something the attempter did not.

The expected values are easy and disconcerting. A ten-year sentence delivered with probability one in 10,057 has an expected size of 8.7 hours. That is the full deterrent weight, per drive, of the entire homicide branch of the law as it bears on a drunk driver. The arrest branch fires about sixty-five times more often, and carries nothing remotely like a decade. Whatever is keeping people from driving drunk, it is not the arithmetic of expected punishment.

The drafters of the American Model Penal Code took the clean position. Section 5.05(1) grades attempt, solicitation and conspiracy at the same level as the offence attempted, with one exception, for first-degree felonies and capital crimes [16]. That is the control principle written into a statute: identical conduct, identical intent, identical liability, outcome irrelevant. And it is where the argument met the world, because jurisdictions divided on whether to follow — some set attempt a class below the completed offence, others match them. Sanford Kadish concluded that the differential is defensible only as a concession to the sentiments Smith described, not as a proposition about desert [17]; David Enoch finds the same split running through every corner of tort and criminal doctrine [18].

## Why It Never Feels Like a Lottery

---

There is a last empirical wrinkle, and it explains why the puzzle does not dissolve on contact with the argument. Markus Kneer and Edouard Machery ran moral-luck pairs past subjects in two designs. Shown both cases side by side, where comparison is unavoidable, the vast majority judged the lucky and unlucky agents equally blameworthy and their actions equally wrong. Shown one case only, outcome moved every judgement measured — wrongness, blame, permissibility, punishment [19]. Fiery Cushman had found a related split earlier: judgements of wrongness and permissibility track what the agent believed and wanted, while judgements of blame and punishment track the causal chain as well [20].

Read those two results together and the shape of the thing appears. Reflective moral judgement is on the side of the control principle. Spontaneous moral judgement is not. And the world is a between-subjects design. The driver who got home is never in the dock beside the driver who did not; nobody convenes the 124 million comparison cases. We meet the unlucky one alone, with the outcome already known, which is the exact condition under which hindsight makes the foreseeable look inevitable and the reckless look monstrous.

Smith's defence survives that, and is strengthened by it. A rule of blame that fires on outcomes is administrable, hard to corrupt, and genuinely informative. What it cannot do is what we think it does: at a correlation of 0.013 it does not sort the dangerous from the safe at all. It picks one member of a class more or less at random and makes an example. Sometimes that is exactly what a legal system needs to do, and the honest version of the defence says so — this is a lottery, we are running it deliberately, and the prize is that no one has to read anybody's mind.

What remains is the residue, and it is not an error to be argued away. The driver whose sidewalk was empty knows perfectly well what he did, and knows equally well that the reason he is not a killer has nothing to do with him. Williams thought the discomfort that produces was not a sign that morality had gone wrong but a sign of what morality had promised and could not deliver — a shelter from luck, built on ground where no such shelter can stand.

---

## REFERENCES

- [1] Centers for Disease Control and Prevention, "Impaired Driving Facts," National Center for Injury Prevention and Control, Atlanta, GA, 2024.
- [2] A. Jewett, R. A. Shults, T. Banerjee, and G. Bergen, "Alcohol-Impaired Driving Among Adults — United States, 2012," *Morbidity and Mortality Weekly Report*, vol. 64, no. 30, pp. 814-817, 2015.
- [3] Centers for Disease Control and Prevention, "Vital Signs: Alcohol-Impaired Driving Among Adults — United States, 2010," *Morbidity and Mortality Weekly Report*, vol. 60, no. 39, pp. 1351-1356, 2011.
- [4] National Highway Traffic Safety Administration, "Alcohol-Impaired Driving: 2023 Data," *Traffic Safety Facts*, Rep. DOT HS 813 713, Washington, DC, 2025.
- [5] Federal Bureau of Investigation, *Crime in the Nation, 2024*, Uniform Crime Reporting Program, Washington, DC, 2025.
- [6] B. Williams, "Moral Luck," *Proceedings of the Aristotelian Society, Supplementary Volumes*, vol. 50, pp. 115-135, 1976.
- [7] T. Nagel, "Moral Luck," *Proceedings of the Aristotelian Society, Supplementary Volumes*, vol. 50, pp. 137-151, 1976.
- [8] B. Williams, "Postscript," in *Moral Luck*, D. Statman, Ed. Albany, NY: State Univ. of New York Press, 1993.
- [9] I. Kant, *Groundwork of the Metaphysics of Morals*, First Section. Riga: J. F. Hartknoch, 1785.
- [10] A. Smith, *The Theory of Moral Sentiments*, Part II, Sect. III. London: A. Millar, 1759.
- [11] M. J. Zimmerman, "Taking Luck Seriously," *Journal of Philosophy*, vol. 99, no. 11, pp. 553-576, 2002.
- [12] J. J. Thomson, "Morality and Bad Luck," *Metaphilosophy*, vol. 20, no. 3-4, pp. 203-221, 1989.
- [13] R. J. Hartman, *In Defense of Moral Luck: Why Luck Often Affects Praiseworthiness and Blameworthiness*. New York: Routledge, 2017.
- [14] Sentencing Council for England and Wales, *Attempted Murder: Definitive Guideline*. London, UK, 2021.
- [15] D. Lewis, "The Punishment That Leaves Something to Chance," *Philosophy & Public Affairs*, vol. 18, no. 1, pp. 53-67, 1989.
- [16] American Law Institute, *Model Penal Code*, §5.05(1). Philadelphia, PA, 1962.

- [17] S. H. Kadish, "Foreword: The Criminal Law and the Luck of the Draw," *Journal of Criminal Law and Criminology*, vol. 84, no. 4, pp. 679-702, 1994.
- [18] D. Enoch, "Moral Luck and the Law," *Philosophy Compass*, vol. 5, no. 1, pp. 42-54, 2010.
- [19] M. Kneer and E. Machery, "No Luck for Moral Luck," *Cognition*, vol. 182, pp. 331-348, 2019.
- [20] F. Cushman, "Crime and Punishment: Distinguishing the Roles of Causal and Intentional Analyses in Moral Judgment," *Cognition*, vol. 108, no. 2, pp. 353-380, 2008.

---

Daily Learning No. 023 · Philosophy · Monday, September 14th, 2026 · Sources are listed above; every load-bearing figure traces to a numbered reference.